

BROCAN

A microservice-based CI



AGENDA

Preliminaries

CI system and messaging primer.

Architecture and Implementation

How Brocan is built up, and how the components fit together.

What happens when...

The way a build is carried out by Brocan.



PRELIMINARIES

PRELIMINARIES

topics

- What is continuous integration?
- Components of a CI system.
- Why use messaging and why use NATS?
- Message pattern matching.

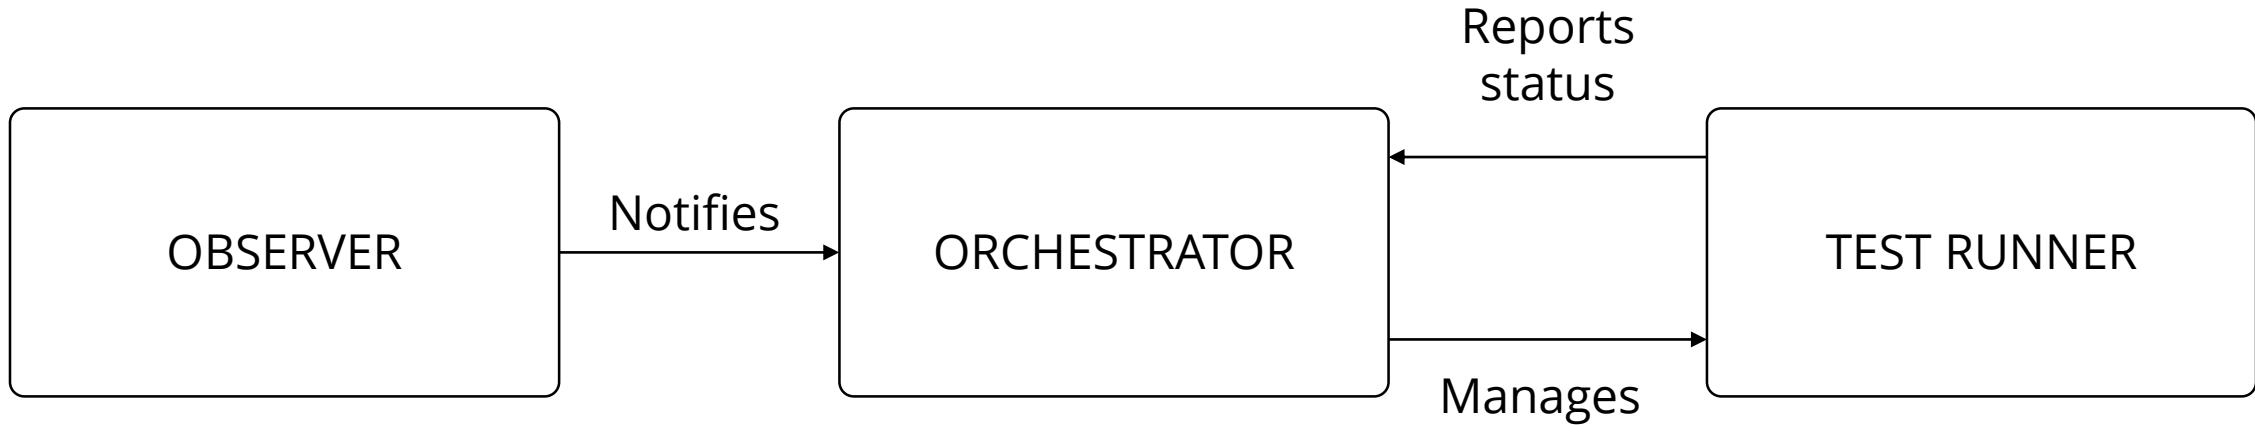
CONTINUOUS INTEGRATION

Continuous Integration is a software development practice where members of a team integrate their work frequently, usually each person integrates at least daily - leading to multiple integrations per day. Each integration is verified by an automated build (including test) to detect integration errors as quickly as possible.

Martin Fowler

TRADITIONAL CI ARCHITECTURE

- A traditional CI system contains three core parts:
 - Observer,
 - Orchestrator (or Dispatcher),
 - Test Runner (or Agent).



TRADITIONAL CI ARCHITECTURE

observer

- Checks a specified repository for changes.
- Notifies the Orchestrator upon a change.
- Does not care about
 - the way builds are carried out,
 - the result of the builds,
 - anything that involves more than the concept of code/changeset/repository.

TRADITIONAL CI ARCHITECTURE

orchestrator

- Manages the swarm of Test Runners.
- Handles incoming Observer notifications, queues if necessary.
- Dispatches build requests to Test Runners.
- Propagates build results to other services if necessary.
- Does not care about the way builds are performed.

TRADITIONAL CI ARCHITECTURE

test runner

- Does the heavy lifting of the build process.
- Receives build requests from the Orchestrator.
- Reports build results to the Orchestrator.
- Needs some metadata to figure out how testing/building can be done.

MESSAGING

core ideas

- Message-first approach.
 - Messages are more important than the services themselves.
 - The actual sender and receiver do not matter, only the communication itself.
- A service is defined by its incoming and outgoing messages.
 - The way these messages are transported does not matter.
- Can be done with HTTP but works better with custom protocols.

MESSAGING routing and discovery

- How to find out where to deliver the messages?
- Many possible options with different trade-offs
 - Predefined IP address of single services.
 - Predefined domain name of single services.
 - Service discovery of single services.
 - Messages brokers (with service discovery or predefined IP/domain).
 - Gossip- or infection-style protocols.

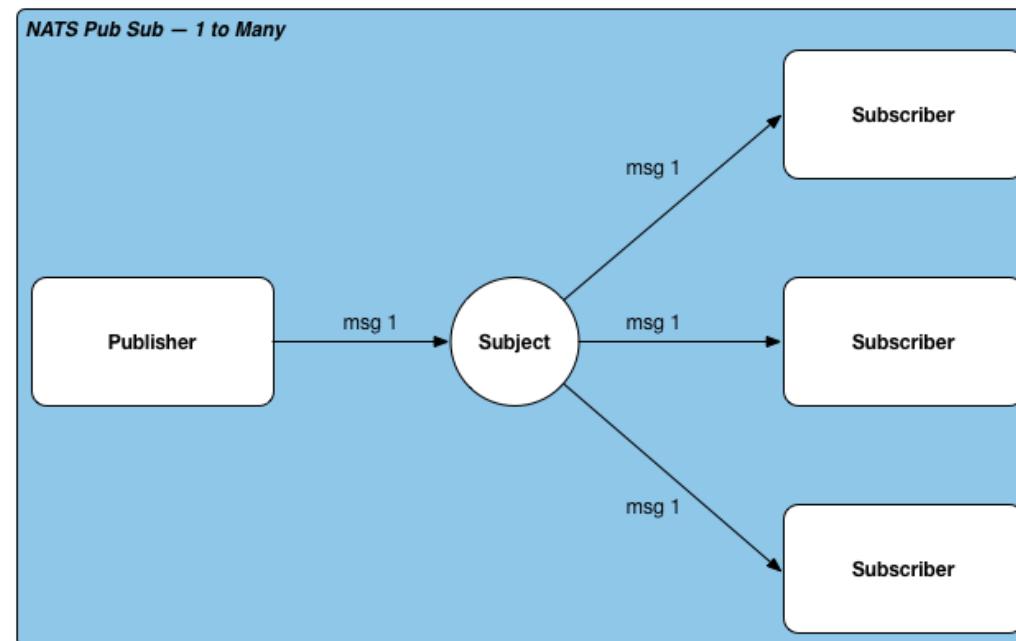
MESSAGING WITH NATS

- NATS is a messaging system managed by a cluster of message brokers.
- High performance (much higher throughput than other brokers).
- Message persistence is not supported out-of-the-box.
- Clients can send messages and subscribe to subjects..
- Various messaging models are supported.

NATS MESSAGING MODELS

publish-subscribe

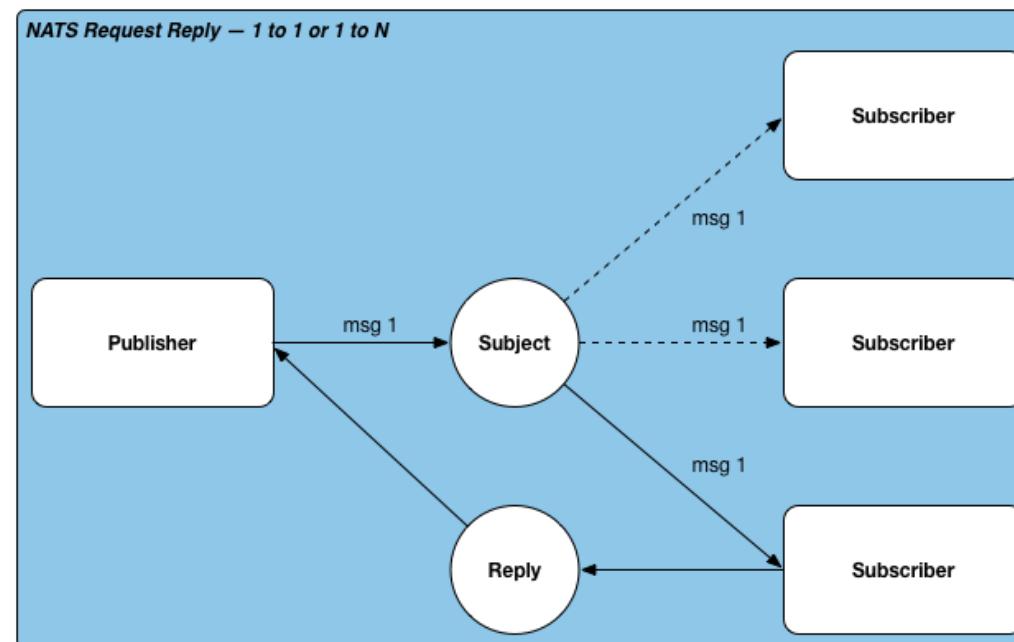
- One-to-many communication.
- Each subscriber of a specific subject receives the message.



NATS MESSAGING MODELS

request-reply

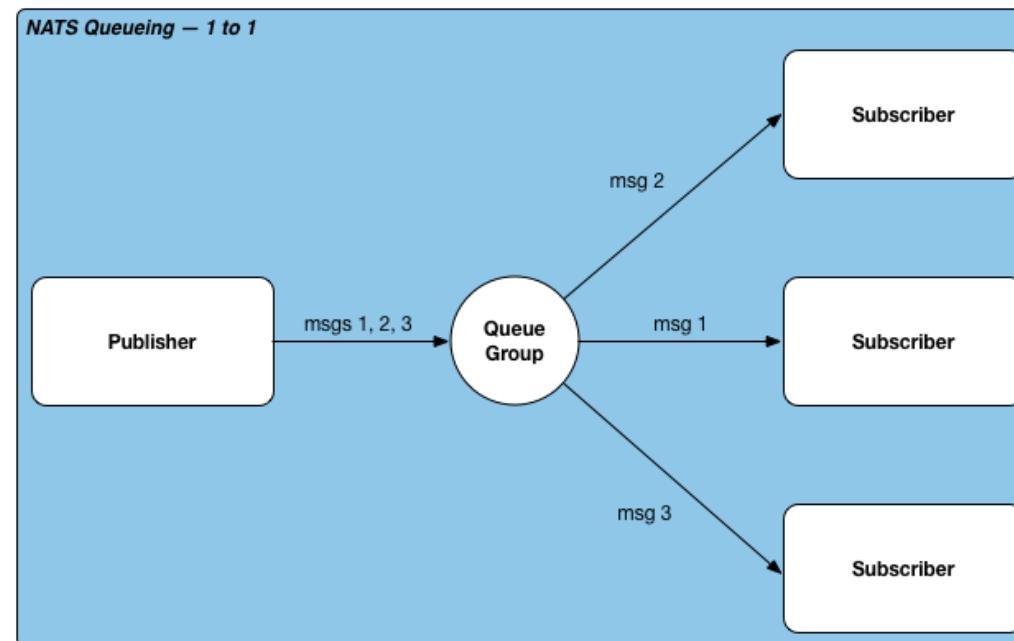
- One-to-one or one-to-many communication.
- The first (or first N) response is returned.



NATS MESSAGING MODELS

queueing

- Subscribers form a queue group.
- Messages are load-balanced between the subscribers.



PATTERN MATCHING

core ideas

- Services can supply a pattern they expect.
 - Messages will be matched against subscriber patterns.
 - Only matching messages will be delivered.
- There are no endpoints anymore – only message patterns.
- Pattern matching makes it easy to incrementally introduce new functionality, to mix and change behaviour.

PATTERN MATCHING

example

- Let's say, we want to calculate the VAT for a shopping cart.
- Different categories of items have different VAT calculation rules.

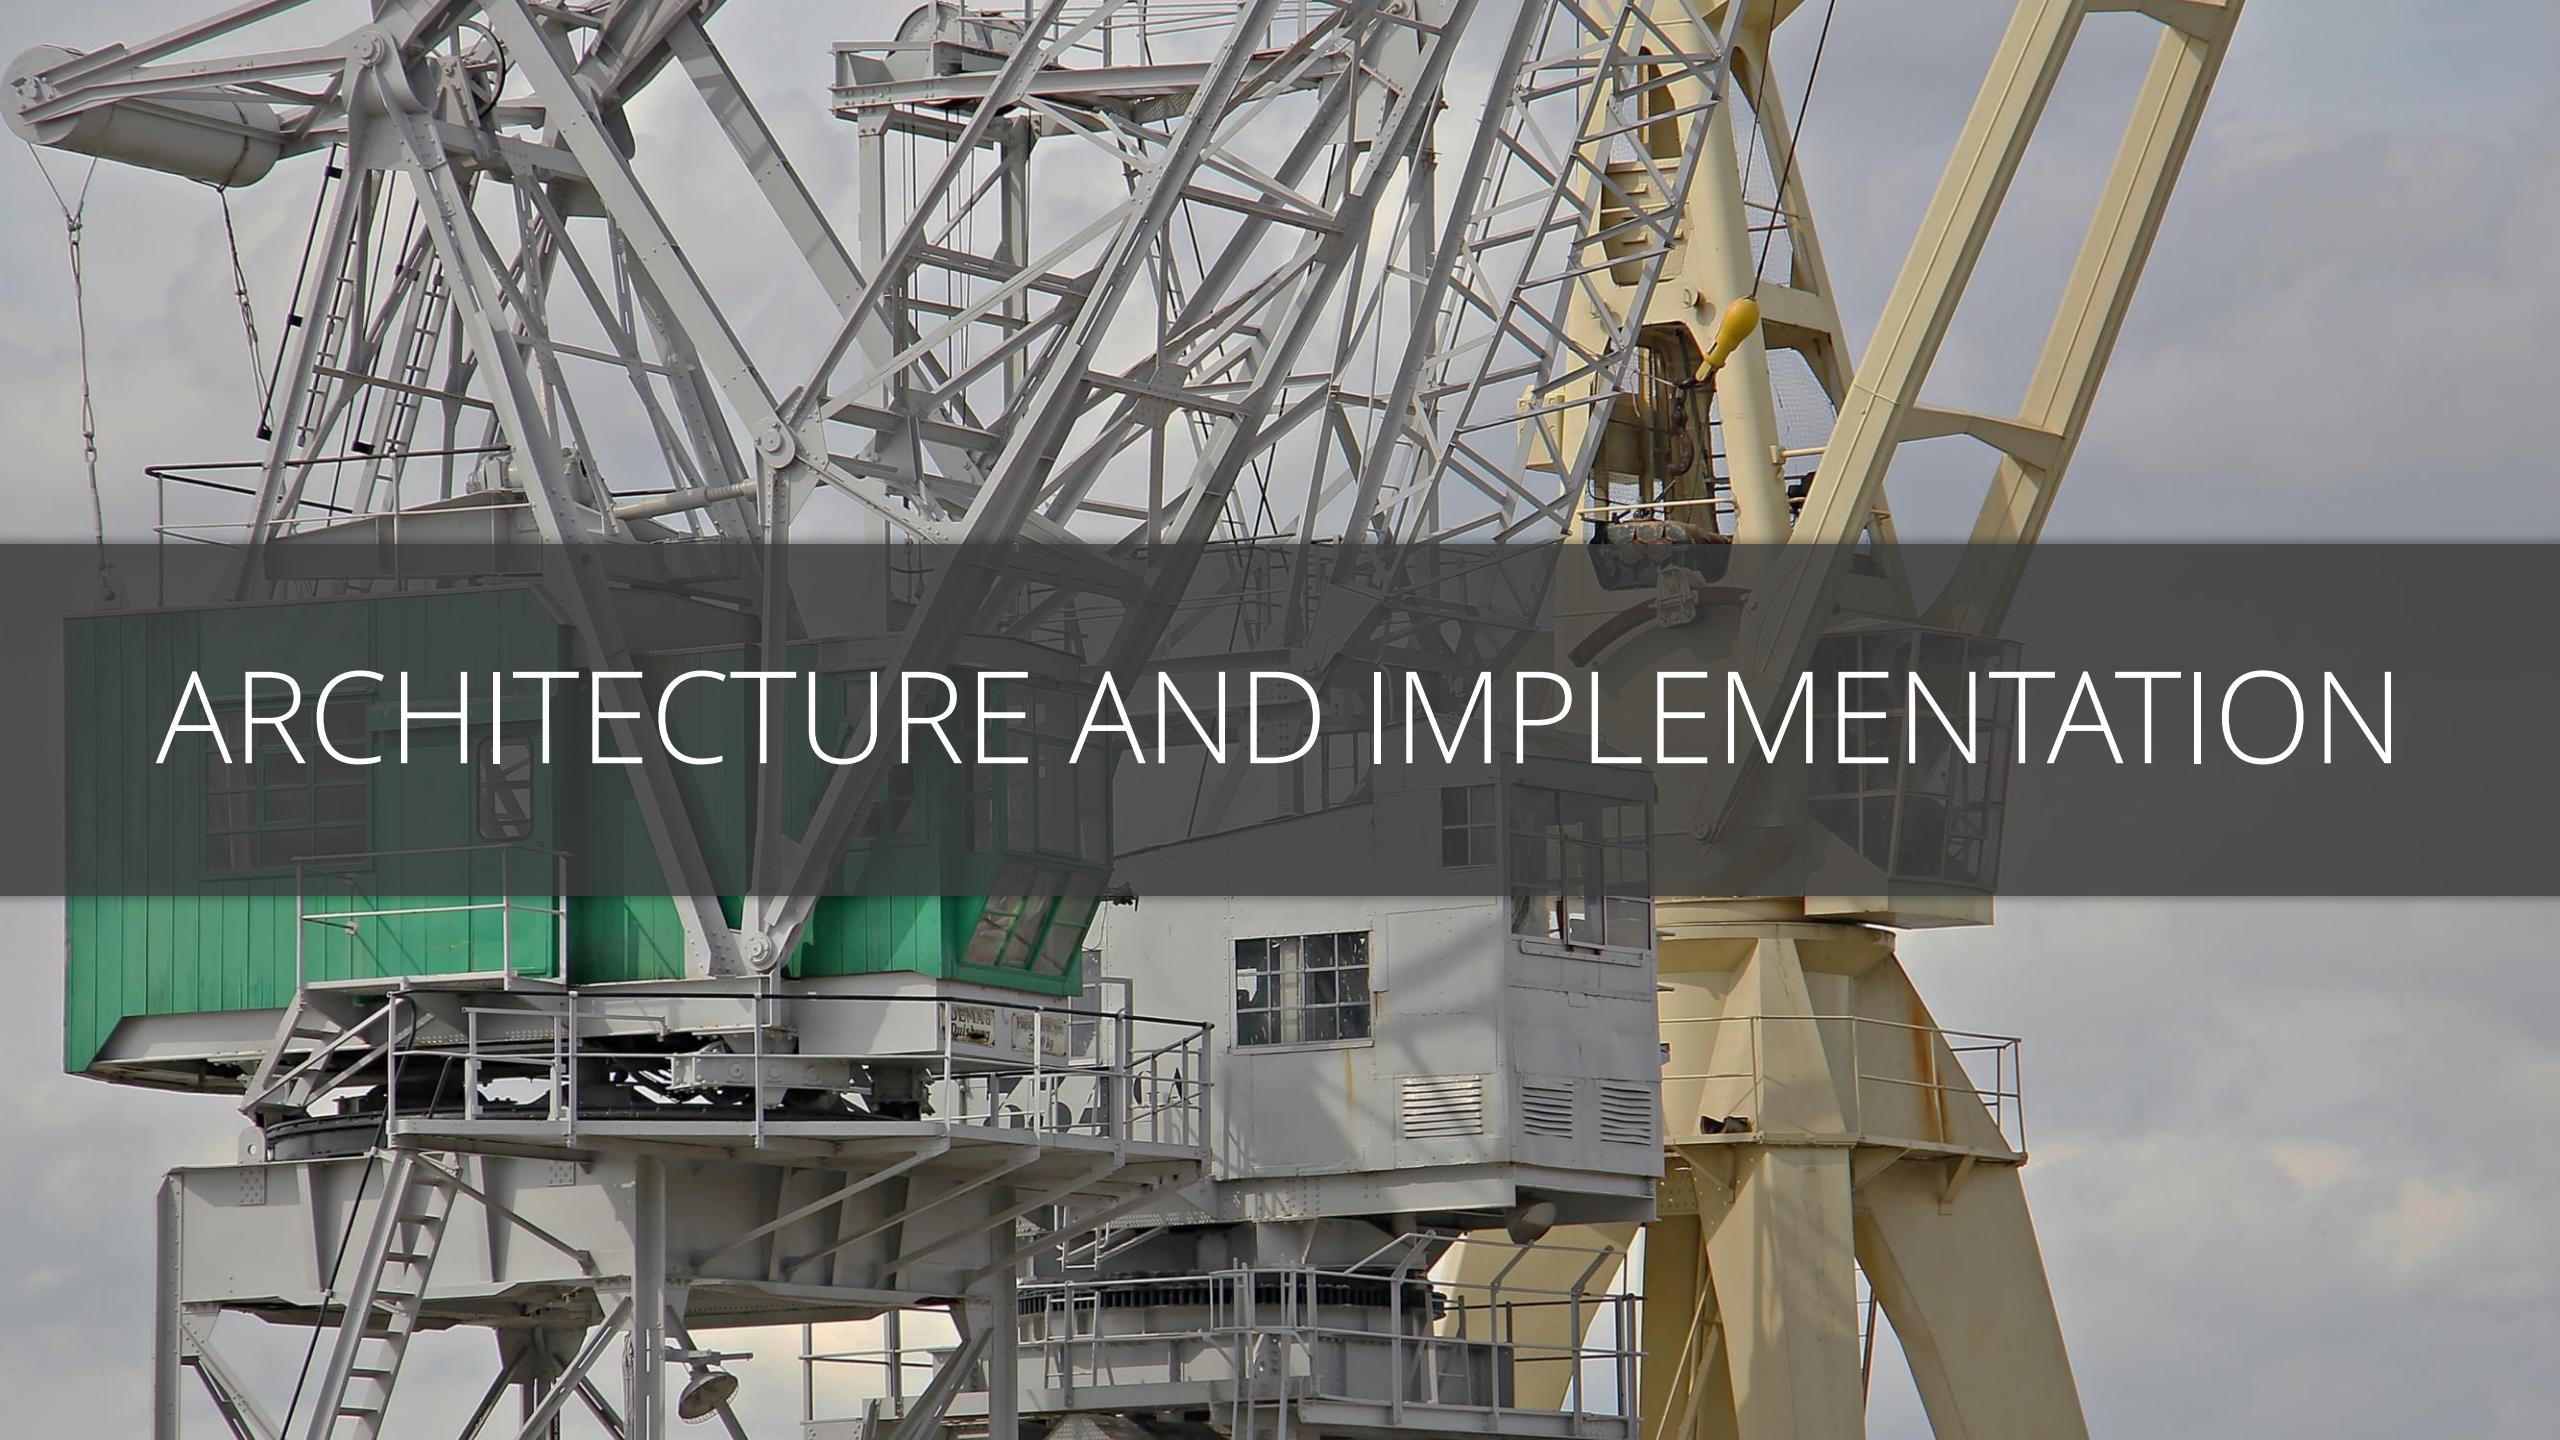
PATTERN MATCHING example

- Imagine, that we publish the following messages:

```
{                                     {  
  name: The Tao Of Microservices    name: Birra Moretti  
  category: [ non-food, book, IT ]  category: [ food, drink, alcoholic, beer ]  
  ...                                ...  
}                                     }
```

- Then the following patterns can be introduced:

```
{                               {           {  
  category: book      category: beer      // catch-all  
}       }           }  
}
```

The background image shows a massive industrial lattice-boom crane, likely a port crane, against a backdrop of a cloudy sky. The crane's structure is made of white-painted steel beams, with a yellow lattice boom extending to the right. A green shipping container is visible behind the crane's body. The overall scene conveys a sense of heavy industry and engineering.

ARCHITECTURE AND IMPLEMENTATION

ARCHITECTURE AND IMPLEMENTATION topics

- Inverted CI architecture.
- Implementation challenges.
- Pattern matching in action.

INVERTED CI ARCHITECTURE

- Externalized Observer.
 - Notifications are supplied by repository providers (eg. GitHub).
- The absence of a central Orchestrator, that's the inversion.
 - Test Runners are not given builds, but take them.
 - The Orchestrator is reduced into a build queue.

ARCHITECTURE ANGLES

message-first

- There's no direct dependency between services.
 - They are just producers and consumers of messages.
- Processes are formed by message flows.
- In Brocan, the only services that have knowledge of each other are the Test Runner and the Agent services.
 - They utilize a JSON HTTP API.

ARCHITECTURE ANGLES

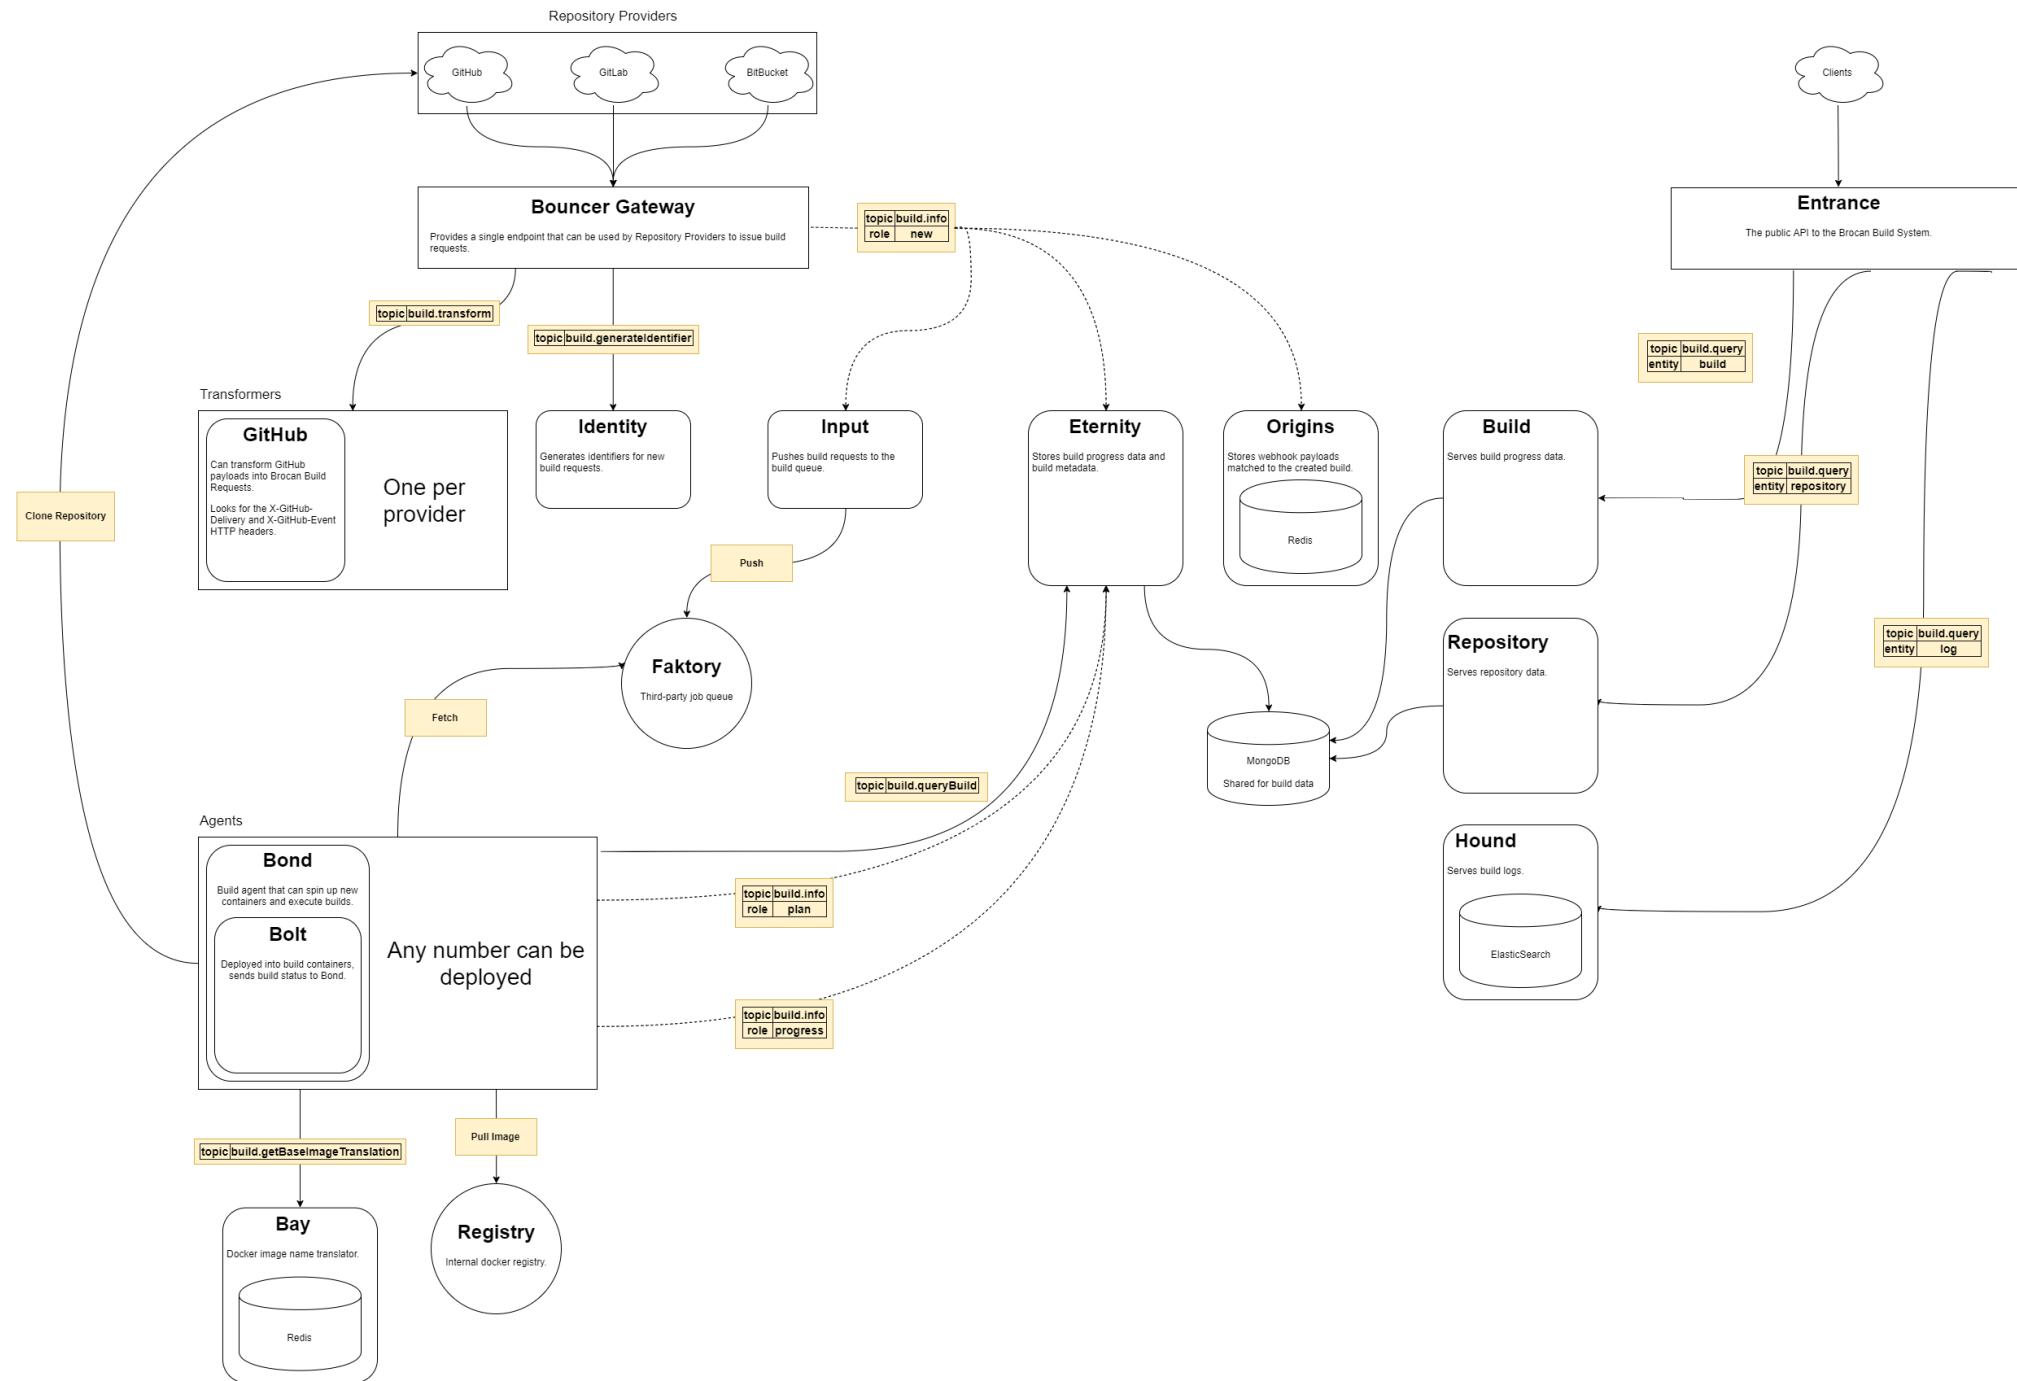
disposable

- Inspiration from Richard Rodger's book, The Tao of Microsevices.
- Services are designed to be disposable.
- Every service is
 - small,
 - concentrated,
 - and atomic.
- The largest service is 805 LOC!

ARCHITECTURE ANGLES

evolution and dynamism

- Being message-first and disposable embraces that individual services are not important.
- This leads to the importance of the processes and flows.
- As business requirements and processes change
 - the system can be easily adjusted and thus evolve,
 - services can be thrown out, merged and introduced on demand.
- In conclusion, the system becomes a living thing, with the messaging system as its nervous system.



IMPLEMENTATION CHALLENGES

achieving distilled business logic

- Distilled business logic refers to services that
 - are only concerned with business tasks,
 - do not care about circuit breaking, service discovery, log transportation etc.
- Nearly impossible to achieve but worth to strive for.
- Resolution for Brocan:
 - Passed with issues, as some non-business code is contained within each service.

IMPLEMENTATION CHALLENGES

multilingualism

- A system should not be framework- or language-locked
 - That would break the „Use right tool for the right job.” principle.
- Regarding the transport, Brocan is locked to NATS.
 - This is not an issue, since NATS has client libraries in multiple languages.
- However, pattern matching and other advanced features are provided by HemeraJS.
 - Therefore Brocan is locked to HemeraJS.

PATTERN MATCHING IN ACTION

webhook transformers

- Repository providers (such as GitHub) can be integrated with transformers.
- Transformers produce the internal BBRF format.
- New transformers can be added with zero infrastructure change thanks to pattern matching.

PATTERN MATCHING IN ACTION

webhook transformers

- Currently only GitHub is integrated, using the following pattern:

```
{  
  webhookRequest: {  
    headers: {  
      x-github-delivery: /.*/,  
      x-github-event: /.*/  
    }  
  }  
}
```

TALKING ABOUT NUMBERS...

13
SERVICES

23
CONTAINERS

8
TOPICS

WHAT HAPPENS WHEN...



WHAT HAPPENS WHEN... topics

- How a build is performed.
- The brocanfile format.
- What gets stored in the DB.



BUILD PROCESS

- It can be seen as two distinct processes:
 - Build Request – Processing and putting a notif into the queue.
 - Build Execution – Actually performing a queued build.
- Prerequisites:
 - The targeted repository contains a brocanfile.
 - A proper webhook is set up for push events.

BUILD PROCESS

brocanfile

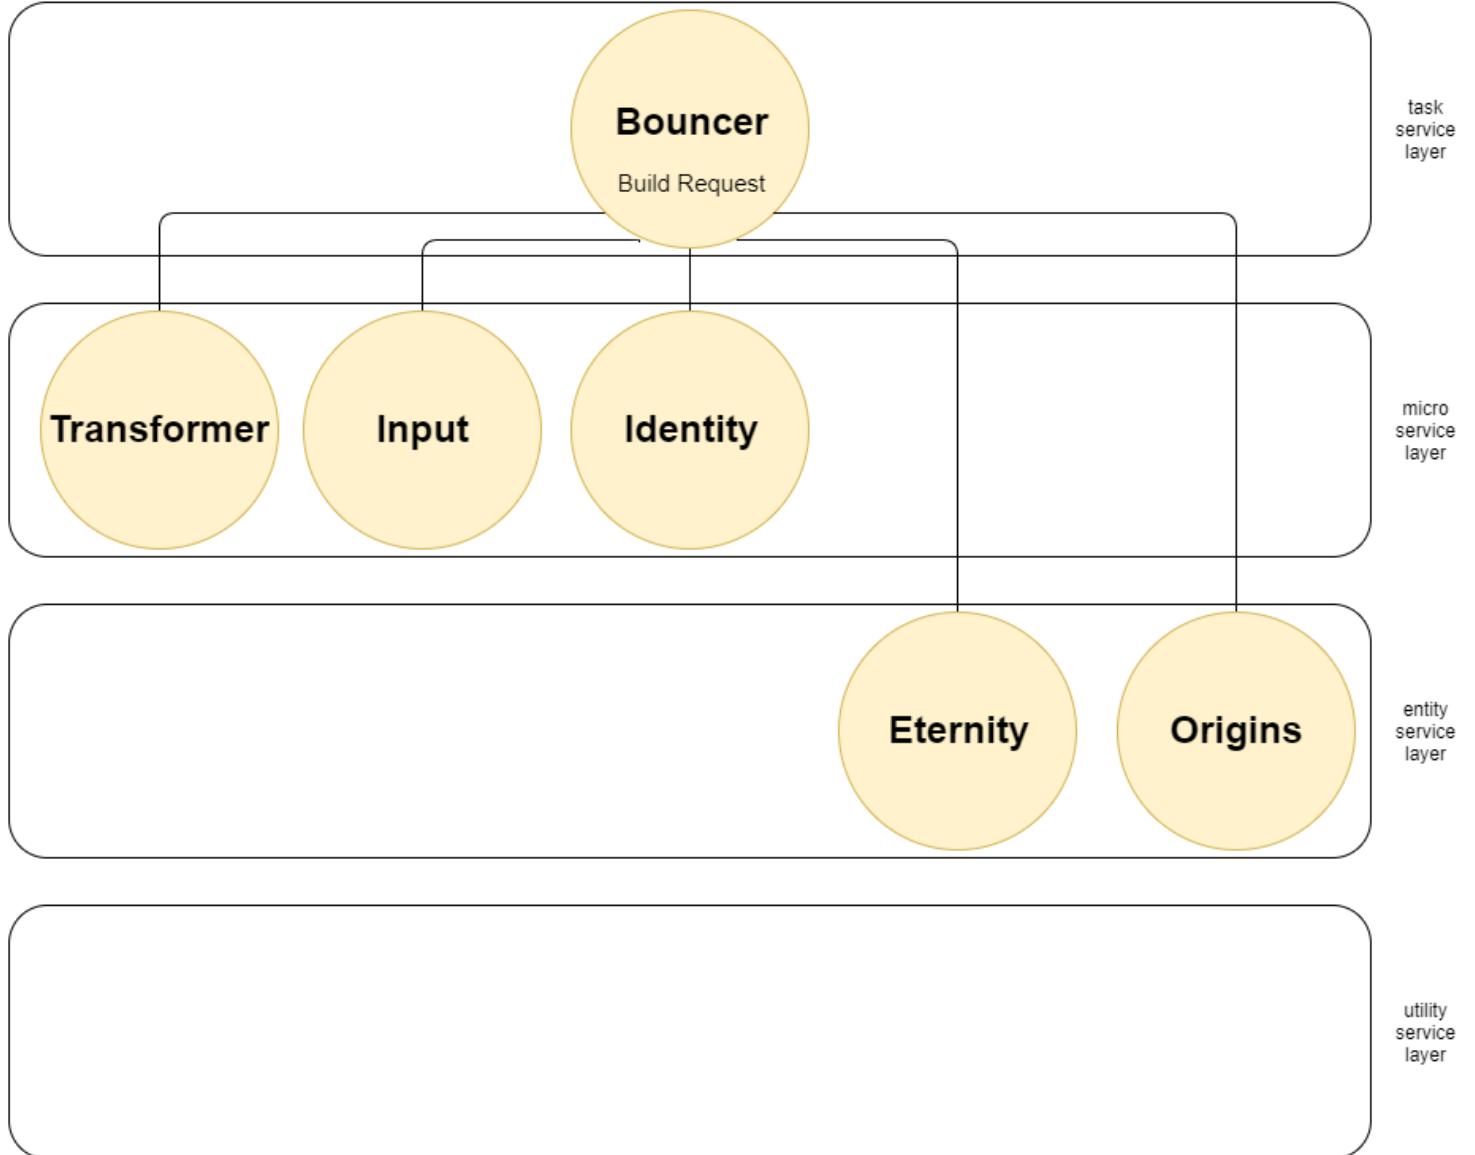
- Based on the HJSON format.
- Properties
 - Base: The environment to run the build in.
 - Owner: The owner of the project.
 - Steps: Step definitions with commands.

```
{  
  base: maven  
  owner: battila7  
  steps: [  
    {  
      name: compile  
      commands: [  
        echo start  
        mvn clean compile  
      ]  
    }  
    {  
      name: test  
      commands: [  
        mvn test  
      ]  
    }  
  ]  
}
```

BUILD PROCESS request

1. New notification received from a repository provider.
2. Notification is transformed into an internal BBRF object.
3. A new Build Identifier is generated for the BBRF.
4. The build is stored and is put into the build queue.

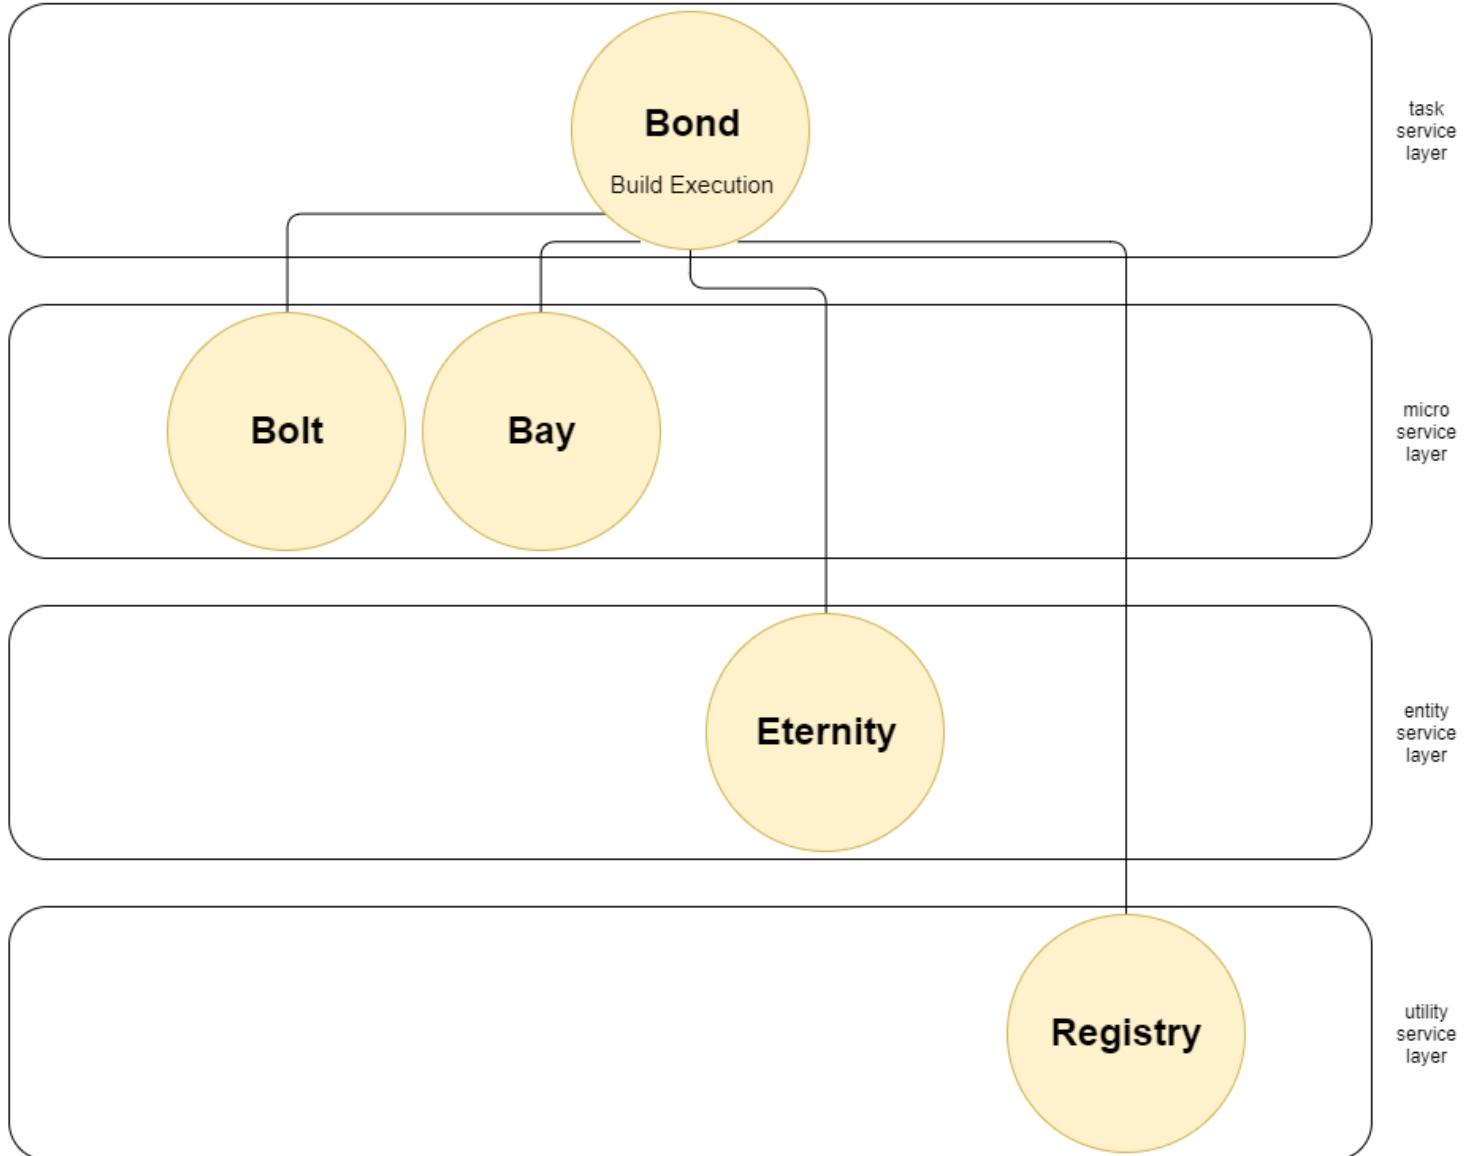
BUILD PROCESS request



BUILD PROCESS execution

1. Once an agent frees up, it queries a build from the queue.
2. The appropriate repository (at the state of the changes) is cloned.
3. A new build container is started, with the runner deployed.
4. As the build container exits, results are propagated into the database.

BUILD PROCESS execution



MONGODB document structure

_id	Calculated from the commit hash, timestamp and branch name.
author	The user/entity who has pushed the changes.
branch	The branch the changes reside on.
commit	Head commit of the changes.
repository	Data related to the containing repository of the changes.
execution	Results of the step and command executions.

MONGODB example document

```
{  
    "_id" : "1abb95fd469c67969e8f6a02c705a71c0ea0e1ab4410625338e94a93383935ca",  
    "timestamp" : 1511208107436,  
    "author" : {  
        "name" : "battila7",  
        "username" : "battila7",  
        "uri" : "https://github.com/battila7"  
    },  
    "branch" : {  
        "name" : "master",  
        "uri" : "https://github.com/battila7/cd2t-100/tree/master"  
    },  
    "commit" : {  
        "hash" : "30d7dca8936847b83c2e908b3f0c2e38473866a7",  
        "message" : "Updated the brocanfile to test Brocan",  
        "uri" : "https://github.com/battila7/cd2t-100/commit/30d7dca8936847b83c2e908b3f0c2e38473866a7"  
    },  
    "repository" : {  
        "name" : "cd2t-100",  
        "uri" : "https://github.com/battila7/cd2t-100"  
    },  
    "execution" : {  
        "steps" : [  
            {  
                "name" : "compile",  
                "commands" : [  
                    {  
                        "command" : "echo start",  
                        "status" : "successful"  
                    },  
                    {  
                        "command" : "mvn clean compile",  
                        "status" : "successful"  
                    }  
                ],  
                "status" : "successful"  
            },  
            {  
                "name" : "test",  
                "commands" : [  
                    {  
                        "command" : "mvn test",  
                        "status" : "failed"  
                    }  
                ],  
                "status" : "failed"  
            }  
        ],  
        "status" : "failed"  
    },  
    "startedAt" : 1511208127653,  
    "finishedAt" : 1511208194076  
}
```



DEMONSTRATION



Q&A



A close-up photograph of an eagle's head, showing its sharp yellow beak and intense yellow eye with a dark pupil. The feathers on the head are white with brown streaks. A solid black rectangular box is overlaid across the middle of the image, containing the text "THANKS FOR YOUR ATTENTION!" in a large, white, sans-serif font.

THANKS FOR YOUR ATTENTION!